



Harvard University SIMR (Summer Introduction to Mathematical Research) **Dinic's Algorithm**

Yongwhan Lim
Research Mentor in SIMR at Harvard
Lecturer in EECS at MIT
Associate in Computer Science at Columbia
7:45pm ET, Thursday, July 7, 2022



Yongwhan Lim



- Currently:
 - (July 11~) Senior Quantitative Software Engineer at Two Sigma
 - Research Mentor in SIMR at Harvard University
 - Lecturer in EECS at MIT
 - Associate in Computer Science at Columbia University
 - ICPC Head Coach at Columbia University
 - ICPC Judge for Mid-central and Greater New York regions in N.A.
- Previously:
 - Research Software Engineer at Google Research
- Education:
 - Stanford: Math & CS (BS '11) and CS (MS '13)
 - MIT: Operations Research (PhD, started in 2013 but on an extended leave-of-absence since 2016)



Lecture Overview on Dinic's Algorithm

- Definitions
- **Dinic's Algorithm**
- Number of Phases
- Proof of Correctness
- Finding Blocking Flow
- Complexity
- Unit Networks
- Unit Capacity Networks
- Implementations in C++
- Looking ahead!
- References
- Contact Information

Lecture Overview on Dinic's Algorithm

- Definitions
- **Dinic's Algorithm**
- Number of Phases
- Proof of Correctness
- Finding Blocking Flow
- Complexity
- Unit Networks
- Unit Capacity Networks
- **Implementations in C++**
- Looking ahead!
- References
- Contact Information

Definitions

- A **residual network** G^R of network G is a network which contains two edges for each edge $(v, u) \in G$:
 - (v, u) with capacity $c_{vu}^R = c_{vu} - f_{vu}$
 - (u, v) with capacity $c_{uv}^R = f_{vu}$
- A **blocking flow** of some network is such a flow that every path from s to t contains at least one edge which is saturated by this flow.
 - Note that a blocking flow is not necessarily maximal.
- A **layered network** of a network G is a network built in the following way:
 - For each vertex v we calculate $\text{level}[v]$: the shortest path (unweighted) from s to this vertex using only edges with positive capacity.
 - We keep only those edges (v, u) for which $\text{level}[v]+1 = \text{level}[u]$.
 - Obviously, this network is acyclic.

Dinic's Algorithm

- The algorithm consists of several phases.
- On each phase:
 - Construct the layered network of the residual network of G .
 - Find an arbitrary blocking flow in the layered network and add it to the current flow.

Number of Phases

- The algorithm terminates in less than V phases.
- Let's prove it using following two lemmas:
 - Lemma 1. The distances from s to each vertex do not decrease after each iteration: that is, $\text{level}_{i+1}[v] \geq \text{level}_i[v]$.
 - Lemma 2. $\text{level}_{i+1}[t] > \text{level}_i[t]$.
- There are less than V phases because $\text{level}[t]$ increases, but it can't be greater than $V - 1$.

Lemma 1: Proof

- $\text{level}_{i+1}[v] \geq \text{level}_i[v]$.
- Fix a phase i and a vertex v . Consider any shortest path P from s to t in G_{i+1}^R . The length of P equals $\text{level}_{i+1}[v]$. Note that G_{i+1}^R can only contain edges from G_i^R and back edges for edges from G_i^R .
- If P has no back edges for G_i^R then $\text{level}_{i+1}[v] \geq \text{level}_i[v]$ because P is also a path in G_i^R . Suppose P has at least one back edge, say, (u, w) . Then, $\text{level}_{i+1}[u] \geq \text{level}_i[u]$.
 - $\text{level}_i[u] = \text{level}_i[w] + 1$: Since (u, w) does not belong to G_i^R , (w, u) was affected by the blocking flow on the previous iteration. Also, $\text{level}_{i+1}[w] = \text{level}_{i+1}[u] + 1$ and $\text{level}_{i+1}[u] \geq \text{level}_i[u]$.
- So, $\text{level}_{i+1}[w] = \text{level}_{i+1}[u] + 1 \geq \text{level}_i[u] + 1 = (\text{level}_i[w] + 1) + 1 = \text{level}_i[w] + 2$.
- Similarly for the rest of the path.

Lemma 2: Proof (by contradiction)

- $\text{level}_{i+1}[t] > \text{level}_i[t]$
- From the previous lemma, $\text{level}_{i+1}[t] \geq \text{level}_i[t]$.
- Suppose that $\text{level}_{i+1}[t] = \text{level}_i[v]$.
- Note that G_{i+1}^R can only contain edges from G_i^R and back edges for edges from G_i^R .
- It means that there is a shortest path in G_i^R which was not blocked by the blocking flow.
- This is a contradiction.

Proof of Correctness

- The algorithm terminates in less than V phases.
- Now, when the algorithm terminated, it could not find a blocking flow in the layered network.
- So, the layered network does not have any path from s to t .
- So, the residual network does not have any path from s to t .
- Therefore, the flow has to be maximum.

Finding Blocking Flow

- In order to find the blocking flow on each iteration, we may simply try pushing flow with DFS (Depth-First Search) from s to t in the layered network while it can be pushed.
- In order to do it more efficiently, we must remove the edges which cannot be used to push anymore.
- We can keep a **pointer** in each vertex which points to the next edge which can be used.

Finding Blocking Flow (con't)

- A single DFS run takes $O(k + V)$ time, where k is the number of pointer advances on this run.
- Over all runs, a number of pointer advances cannot exceed E .
- A total number of runs would not exceed E , as every run saturates at least one edge.
- So, a total running time of finding a blocking flow is **$O(VE)$** .

Complexity

- Since there are less than V phases, the total time complexity is $\mathbf{O(V^2E)}$.

Unit Networks

- A **unit network** is a network in which for any vertex except vertex s and vertex t , **either incoming or outgoing edge is unique and has unit capacity**.
 - This is the setting with the network built to solve a maximum matching problem with flows.
- On unit networks, Dinic's algorithm runs in $O(EV^{1/2})$.

Unit Networks: Dinic's algorithm in $O(EV^{1/2})$

- Each phase now works in $O(E)$ because each edge will be considered at most once. So, it suffices to see a number of phases is $O(V^{1/2})$.
- Suppose there have already been $V^{1/2}$ phases.
 - All the augmenting paths with the length $\leq V^{1/2}$ have been found.
 - Let f be the current flow, f' be the maximum flow and consider their difference $f' - f$.
 - It is a flow in G^R of value $\text{val}(f') - \text{val}(f)$ where on each edge it is either 0 or 1.
 - It can be decomposed into $\text{val}(f') - \text{val}(f)$ paths from s to t and possibly cycles.
 - As the network is unit, they cannot share common vertices.
 - Hence, $(\text{val}(f') - \text{val}(f)) V^{1/2} \leq (\text{the total number of vertices}) \leq V$.
- Therefore, in another $V^{1/2}$ iterations, we will definitely find the maximum flow.

Unit Capacity Networks: : Dinic in $O(E^{3/2})$

- In a more general settings where all edges still have unit capacities (but a constraint on *the number of incoming and outgoing edges* is relaxed), the paths cannot have common edges as opposed to common vertices.
- So, a number of phases is $O(E^{1/2})$, leading to the algorithm running in $O(E^{3/2})$.
- Alternatively, a number of phases can be shown to be $O(V^{2/3})$, which gives even better bound on dense graphs.

Implementations in C++: *FlowEdge* struct

```
struct FlowEdge {  
    int v, u;  
    long long cap, flow = 0;  
    FlowEdge(int v, int u, long long cap) : v(v), u(u), cap(cap) {}  
};
```

Implementations in C++: *Dinic* struct

```
struct Dinic {
    const long long flow_inf = 1e18;
    vector<FlowEdge> edges;
    vector<vector<int>> adj;
    int n, m = 0;
    int s, t;
    vector<int> level, ptr;
    queue<int> q;
    Dinic(int n, int s, int t) : n(n), s(s), t(t) {
        adj.resize(n);
        level.resize(n);
        ptr.resize(n);
    }
    void add_edge(int v, int u, long long cap);
    bool bfs();
    long long dfs(int v, long long pushed);
    long long flow();
};
```

Implementations in C++: *flow* function

```
long long flow() {  
    long long f = 0;  
    while (true) {  
        fill(level.begin(), level.end(), -1);  
        level[s] = 0;  
        q.push(s);  
        if (!bfs())  
            break;  
        fill(ptr.begin(), ptr.end(), 0);  
        while (long long pushed = dfs(s, flow_inf)) {  
            f += pushed;  
        }  
    }  
    return f;  
}
```

Implementations in C++: *bfs* function

```
bool bfs() {
    while (!q.empty()) {
        int v = q.front();
        q.pop();
        for (int id : adj[v]) {
            if (edges[id].cap - edges[id].flow < 1)
                continue;
            if (level[edges[id].u] != -1)
                continue;
            level[edges[id].u] = level[v] + 1;
            q.push(edges[id].u);
        }
    }
    return level[t] != -1;
}
```

Implementations in C++: *dfs* function

```
long long dfs(int v, long long pushed) {
    if (pushed == 0)
        return 0;
    if (v == t)
        return pushed;
    for (int& cid = ptr[v]; cid < (int)adj[v].size(); cid++) {
        int id = adj[v][cid];
        int u = edges[id].u;
        if (level[v] + 1 != level[u] || edges[id].cap - edges[id].flow < 1)
            continue;
        long long tr = dfs(u, min(pushed, edges[id].cap - edges[id].flow));
        if (tr == 0)
            continue;
        edges[id].flow += tr;
        edges[id ^ 1].flow -= tr;
        return tr;
    }
    return 0;
}
```

Implementations in C++: *add_edge* function

```
void add_edge(int v, int u, long long cap) {  
    edges.emplace_back(v, u, cap);  
    edges.emplace_back(u, v, 0);  
    adj[v].push_back(m);  
    adj[u].push_back(m + 1);  
    m += 2;  
}
```

Looking ahead!

- Nearly-linear(!) time algorithm (Li Chen, Rasmus Kyng, Yang P. Liu, Richard Peng, Maximilian Probst Gutenberg, Sushant Sachdeva)
 - Video: https://www.youtube.com/watch?v=7ivZ0_vx4Nc
 - Try to watch the video before Professor Richard Peng's lecture next week! :)
- Dynamic graphs! (Professor Richard Peng's lecture, scheduled on July 14)

References

- <https://cp-algorithms.com/graph/dinic.html>
- <https://people.orie.cornell.edu/dpw/orie633/LectureNotes/lecture9.pdf>
- <http://www.cs.cmu.edu/afs/cs/academic/class/15451-f14/www/lectures/lec11/dinic.pdf>
- <http://www.cs.toronto.edu/~bor/375s06/dinic-sketch.pdf>
- Introduction to Algorithms (CLRS), 4th Edition
- Wikipedia

Contact Information

- yongwhan@mit.edu or yongwhan.lim@columbia.edu
 - <https://cs.columbia.edu/~yongwhan>
 - <https://www.linkedin.com/in/yongwhan/>
-
- For those who are interested in **quantitative analysis in finance**, feel free to send me an email to schedule 1:1 zoom meeting.